

Do current reasoning models allocate more thinking to high-consequence tasks?

A 3-model audit on 300 SWE-bench Lite tasks.



Qwen3-8B

consequence sensitivity

1%

$p(\text{thinking}, \text{consequence}) = 0.002$

👤 n = 50

⚙️ max_new = 4096

🔍 Finding: almost no consequence sensitivity



Qwen3-VL-8B-Thinking

saturation rate

99%

99.3% hit the 8192-token cap

👤 n = 50

⚙️ max_new = 8192

🔍 Finding: nearly all tasks saturate



Claude Sonnet 4.5

consequence sensitivity

20%

$p(\text{thinking}, \text{consequence}) = 0.203$

👤 n = 171 p = 0.0078

📈 High vs Low thinking: +20.5%

🛡️ weak but significant consequence sensitivity