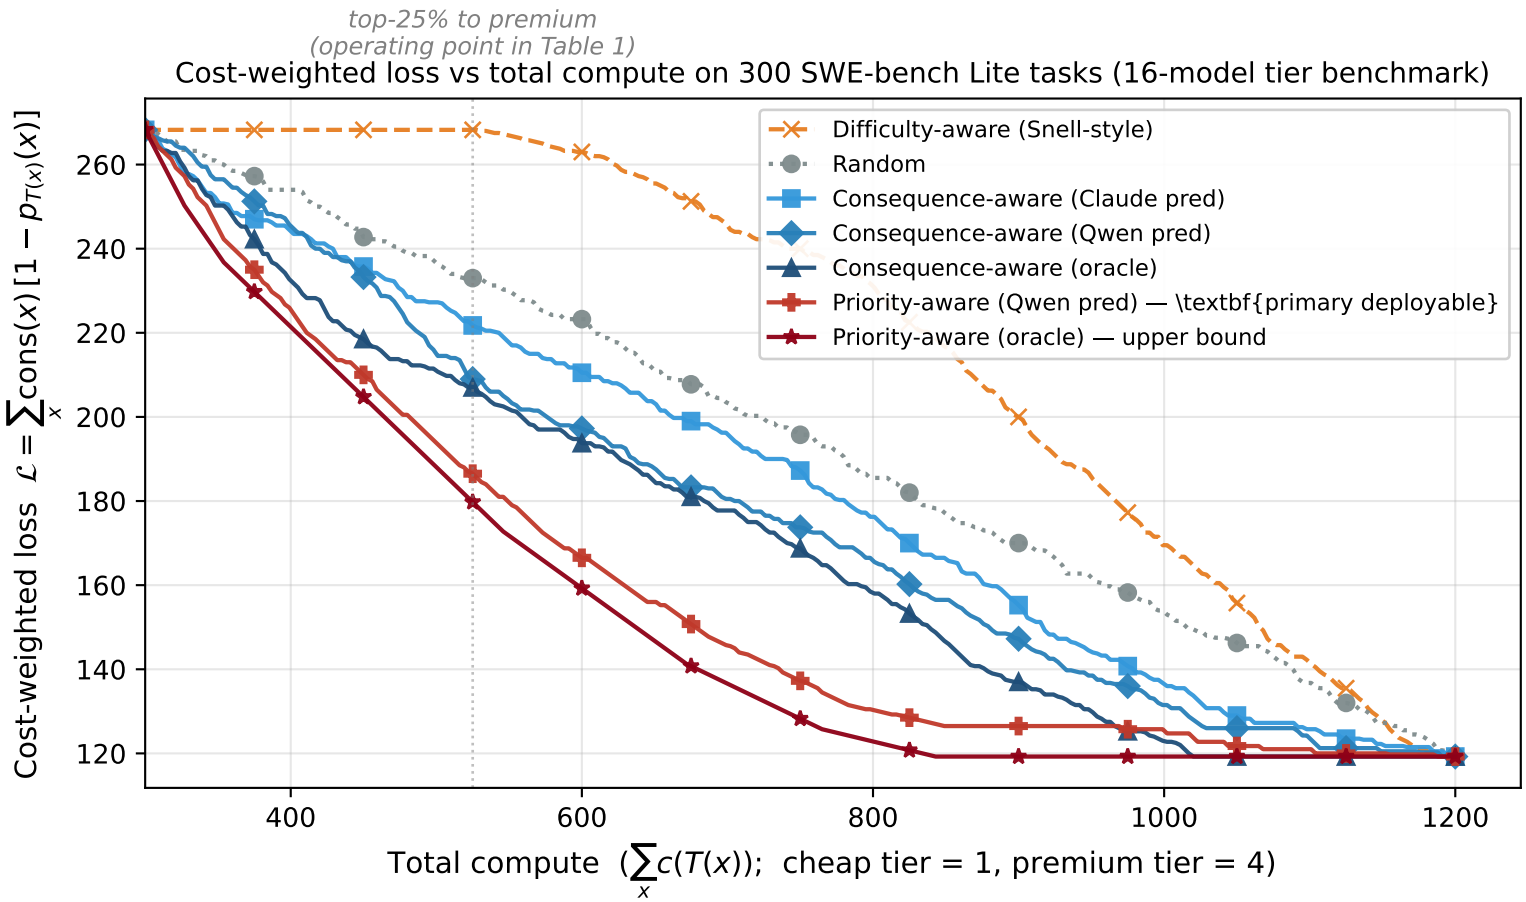
top-25% to premium
(operating point in Table 1)
Cost-weighted loss vs total compute on 300 SWE-bench Lite tasks (16-model tier benchmark)
Difficulty-aware (Snell-style)
Random
Consequence-aware (Claude pred)
Consequence-aware (Qwen pred)
Consequence-aware (oracle)
Priority-aware (Qwen pred) — \textbf{primary deployable}
Priority-aware (oracle) — upper bound
Cost-weighted loss $\mathcal{L} = \sum_x \text{cons}(x)\,[1 - p_{T(x)}(x)]$
Total compute ($\sum_x c(T(x))$; cheap tier = 1, premium tier = 4)
120
140
160
180
200
220
240
260
400
600
800
1000
1200