


Cost-weighted loss vs total compute on 300 SWE-bench Lite tasks (16-model tier benchmark)



- Difficulty-aware (Snell-style)
- Consequence-aware (Qwen pred)
- Consequence-aware (Claude pred)
- Random
- Consequence-aware (oracle)
- Priority-aware (Qwen pred) – primary deployable
- Priority-aware (oracle) – upper bound